

MARKETING AGENT — THE GAP ANALYSIS

How a campaign is made today, and how it changes with the agent

WHAT THIS IS

This is the before-and-after story that justifies the whole project. "Before" is how a marketing campaign gets made inside a human agency. "After" is how the same work happens once a team of specialist AI agents does it. Every feature in the rest of the system traces back to a gap shown here, so this document comes first.

The project brief sets the stakes plainly. Copywriter roles fell 28% in a year, for two years running. One Dutch agency went from 290 copywriters to 20 after adopting AI. The work most exposed is the execution layer — the research, the audience definition, the third draft of someone else's brief. This system automates exactly that layer, so the gap it closes is real and already being closed in the market.


THE WORK TODAY (AS-IS)

A campaign passes from hand to hand. A client sends a brief. An account planner researches the product and the market, then defines a target audience by hand. A copywriter drafts variants, rewrites them after internal feedback, and waits. A legal reviewer reads every line for risky claims and kicks anything doubtful back to the writer. A conversion specialist tunes the call to action for the channel. Then it goes to the client, who asks for revisions, and parts of the loop run again.

Even the smallest agency performs five distinct functions to do this: product and market analysis, audience definition, copy and creative, conversion and call-to-action work, and compliance review. No single person owns the whole chain well — a creative director makes weak media calls, a media buyer writes weak copy. The phases need genuinely different skills, which is why agencies split them across roles in the first place.

The cost of that split is time and hand-offs. The job is measured in days, most of it spent waiting, relaying, and re-reviewing.

AS-IS — How a campaign is made by a human agency



See the picture: [diagrams/gap-as-is.mmd](#)


THE WORK WITH THE AGENT (TO-BE)

The same five functions become five places in one automated pass:

- Product and market analysis -> the vision and color agents read the product image into tags, signals, mood, and a palette.
- Audience definition -> the audience agent turns those signals into one named persona with an age range, interests, and platform behaviour.
- Copy and creative -> the copywriter agent writes two variants for that persona, one emotional and one benefit-led.
- Compliance review -> the guardrails check scans every line for risky claims with plain rules, in well under a second.
- Conversion -> the cta_optimizer agent scores each call to action and suggests a stronger one.

What was days of relay becomes about a minute of machine work. The one human step the "after" picture keeps on purpose is review: a person reads the report and refines it, the same judgement an account manager applies before a client sees the work — and nothing else.

TO-BE — The same five functions as one agent pipeline



See the picture: [diagrams/gap-to-be.mmd](#)

WHY KEEP A HUMAN IN THE LOOP

The brief is blunt about the risk: this kind of system eliminates middle-class jobs, and one agent's mistake can become another agent's input. So the design does not pretend

the output is finished. The campaign is presented as a professional starting point that needs human review, the compliance flags are shown rather than silently fixed, and the refinement panel exists so a person stays in control of tone, audience, and wording. The limits are documented and visible, not hidden.

WHY IT'S BUILT THIS WAY

Because the gap is structural, not cosmetic. The agencies split the work across five roles for a reason — each phase needs different context — so the agents mirror that same split rather than asking one model to do everything at once. The decision to use five specialists instead of one generalist is tested in DL-02; the human-review stance is argued in DL-08. Every later decision log points back to a gap on this page.

WHERE THIS LIVES IN THE CODE

- `app.py` runs the five-function pass in fixed order
- `agents/vision.py` product analysis (image -> tags, signals)
- `agents/color.py` product analysis (image -> palette)
- `agents/audience.py` audience definition (-> persona)
- `agents/copywriter.py` copy and creative (-> two variants)
- `agents/guardrails.py` compliance review (plain-rule claim checks)
- `agents/cta_optimizer.py` conversion (-> CTA scores and suggestions)
- `Project_case.md` the workforce data and the assignment brief

DIAGRAMS FOR THIS DOCUMENT

- `diagrams/gap-as-is.mmd` the manual agency workflow, as swimlanes
- `diagrams/gap-to-be.mmd` the same five functions as one agent pass