

@app — telemetry

Appears in: **DL-08**

app.py:30-46, 175-189

"Append-only JSONL (run_log.jsonl). One line per run ..." ... footnote ¹ — "the logging code is @app, writing to @run-log."

@app — telemetry

app.py:30-46

```
30 def _log_run(output: dict, pipeline_s: float) -> None:
31     cta = output.get("cta_analysis", {})
32     scores = [cta.get(k, {}).get("score", 0) for k in ("variant_1", "variant_2")]
33     valid_scores = [s for s in scores if s]
34     avg_score = round(sum(valid_scores) / len(valid_scores), 1) if valid_scores else 0.0
35     entry = {
36         "ts": time.strftime("%Y-%m-%dT%H:%M:%SZ", time.gmtime()),
37         "run_id": output.get("run_id", ""),
38         "model_used": output.get("vision", {}).get("model_used", ""),
39         "pipeline_s": round(pipeline_s, 1),
40         "compliance_flags": len(output.get("compliance_flags", [])),
41         "cta_score_avg": avg_score,
42         "rag_chunks_used": output.get("rag_chunks_used", 0),
43         "persona": output.get("audience", {}).get("persona_label", ""),
44     }
45     with open(RUN_LOG, "a", encoding="utf-8") as f:
46         f.write(json.dumps(entry) + "\n")
```

...

app.py:175-189

```
175 @app.route("/stats")
176 def stats():
177     if not os.path.exists(RUN_LOG):
178         return jsonify({"runs": 0, "avg_pipeline_s": 0, "avg_cta_score": 0, "compliance_rate": 0})
179     with open(RUN_LOG, encoding="utf-8") as f:
180         lines = [json.loads(ln) for ln in f if ln.strip()]
181     if not lines:
182         return jsonify({"runs": 0, "avg_pipeline_s": 0, "avg_cta_score": 0, "compliance_rate": 0})
183     return jsonify({
184         "runs": len(lines),
185         "avg_pipeline_s": round(sum(l["pipeline_s"] for l in lines) / len(lines), 1),
186         "avg_cta_score": round(sum(l["cta_score_avg"] for l in lines) / len(lines), 1),
187         "compliance_rate": round(sum(1 for l in lines if l["compliance_flags"] > 0) / len(lines), 2),
188         "recent": lines[-5:][::-1],
189     })
```

What it shows: `_log_run` writes ~200 bytes of metrics (no copy text) per run off the critical path; `/stats` aggregates runs, avg pipeline seconds, avg CTA score, and compliance rate.

Source: app.py:30-46, 175-189 · image rendered by `build_snippets.py`