

@config — RAG settings

Appears in: [DL-05](#)

`config.example.py:13-19`

"At 400 characters with 80-character overlap, retrieval pulls coherent sentences; smaller chunks returned mid-sentence fragments." (and "nomic-embed-text via Ollama for embeddings.")

@config — RAG settings

`config.example.py:13-19`

```
13 CHROMA_PERSIST_DIR = "chroma_store"
14 CHROMA_COLLECTION = "marketing_docs"
15 EMBED_MODEL        = "nomic-embed-text"
16 RAG_CHUNK_SIZE     = 400           # characters per chunk
17 RAG_CHUNK_OVERLAP  = 80           # overlap between adjacent chunks
18 RAG_TOP_K           = 4            # chunks retrieved per campaign query
19 DOC_UPLOAD_FOLDER  = "static/doc_uploads"
```

What it shows: The RAG configuration — local Chroma store, nomic embeddings, and the 400/80 chunking that mattered more than the embedding model.

Source: `config.example.py:13-19` · image rendered by `build_snippets.py`