

Stakeholder Analysis

Marketing Agent POC

Primary Stakeholder: Small Business Owner / Solo Marketer

Needs: - Campaign output without hiring an agency or knowing prompt engineering - Two copy variants ready to post, not a rough draft that needs rewriting - Clear indication of any risky claims before publishing - Ability to adjust tone and audience without starting over

Concerns: - Is the output actually usable, or is it obviously AI-written? - Will it hallucinate product specs I didn't provide? - What do I do if the compliance flags are wrong?

Success from their perspective: - Upload a product image, get two variants, post one with minor edits - The persona feels accurate to who actually buys this product - No compliance surprises after publishing

Secondary Stakeholder: Marketing Student / Course Evaluator

Needs: - Visible agent collaboration -- not just output, but evidence of which agents ran - Decision documentation connecting design choices to evidence - A system that demonstrates specialisation rather than just parallel prompts

Concerns: - Is the multi-agent structure genuinely adding value, or just adding complexity? - Are the design decisions documented with real evidence? - Does the system show LO1-LO5 coverage across the full build arc?

Success from their perspective: - The routing label shows which agents ran and why - Each architectural decision has a before/after comparison or a test result - The ethics section grapples with real consequences, not generic platitudes

Tertiary Stakeholder: Copywriter / Creative Professional (Displaced)

Needs: - Transparency about what the system can and can't do - Acknowledgment that this automates real work currently done by people

Concerns: - Will users treat AI copy as finished rather than a starting point? - Is the compliance check comprehensive enough to catch what a human reviewer would? - Who is accountable if the output causes harm?

Success from their perspective: - The system is honest about its limits (persona is a hypothesis, not market research) - Human review is framed as required, not optional - The DL-08 ethics entry names the displacement data directly

Ethical Stakeholder: Future Users and Affected Workforce

Needs: - Documented accountability gaps before they become production incidents - Clear framing of what "professional output" means and doesn't mean

Concerns: - Inter-agent error propagation (hallucinated product tag -> wrong persona -> wrong copy) - Compliance coverage that appears exhaustive but isn't - Persona bias from training data distributions

Success from this perspective: - Three accountability gaps documented in DL-08 - Output framed as a starting point requiring human review, not a finished campaign